

Machine-Level Programming V: Advanced Topics



Today

- **Memory Layout**

- **Buffer Overflow**

 - Vulnerability

 - Protection

- **Unions**



x86-64 Linux Memory Layout

not drawn to scale

00007FFFFFFF

Stack

- Runtime stack (8MB limit)
- E. g., local variables

Heap

- Dynamically allocated as needed
- When call `malloc()`, `calloc()`, `new()`

Data

- Statically allocated data
- E.g., global vars, `static` vars, string constants


Text / Shared Libraries

- Executable machine instructions
- Read-only

Hex Address



400000
000000



Memory Allocation Example


not drawn to scale

```
char big_array[1L<<24]; /* 16 MB */
char huge_array[1L<<31]; /* 2 GB */

int global = 0;

int useless() { return 0; }

int main ()
{
    void *p1, *p2, *p3, *p4;
    int local = 0;
    p1 = malloc(1L << 28); /* 256 MB */
    p2 = malloc(1L << 8); /* 256 B */
    p3 = malloc(1L << 32); /* 4 GB */
    p4 = malloc(1L << 8); /* 256 B */
    /* Some print statements ... */
}
```



Where does everything go?




x86-64 Example Addresses

address range $\sim 2^{47}$

```
local
p1
p3
p4
p2
big_array
huge_array
main()
useless()
```

```
0x00007ffe4d3be87c
0x00007f7262a1e010
0x00007f7162a1d010
0x000000008359d120
0x000000008359d010
0x0000000080601060
0x0000000000601060
0x000000000040060c
0x0000000000400590
```

not drawn to scale



Today

- Memory Layout

- **Buffer Overflow**

 - Vulnerability

 - Protection

- Unions



Recall: Memory Referencing Bug Example

```
typedef struct {
    int a[2];
    double d;
} struct_t;

double fun(int i) {
    volatile struct_t s;
    s.d = 3.14;
    s.a[i] = 1073741824; /* Possibly out of bounds */
    return s.d;
}
```

fun(0)	⌘	3.14
fun(1)	⌘	3.14
fun(2)	⌘	3.1399998664856
fun(3)	⌘	2.00000061035156
fun(4)	⌘	3.14
fun(6)	⌘	Segmentation fault

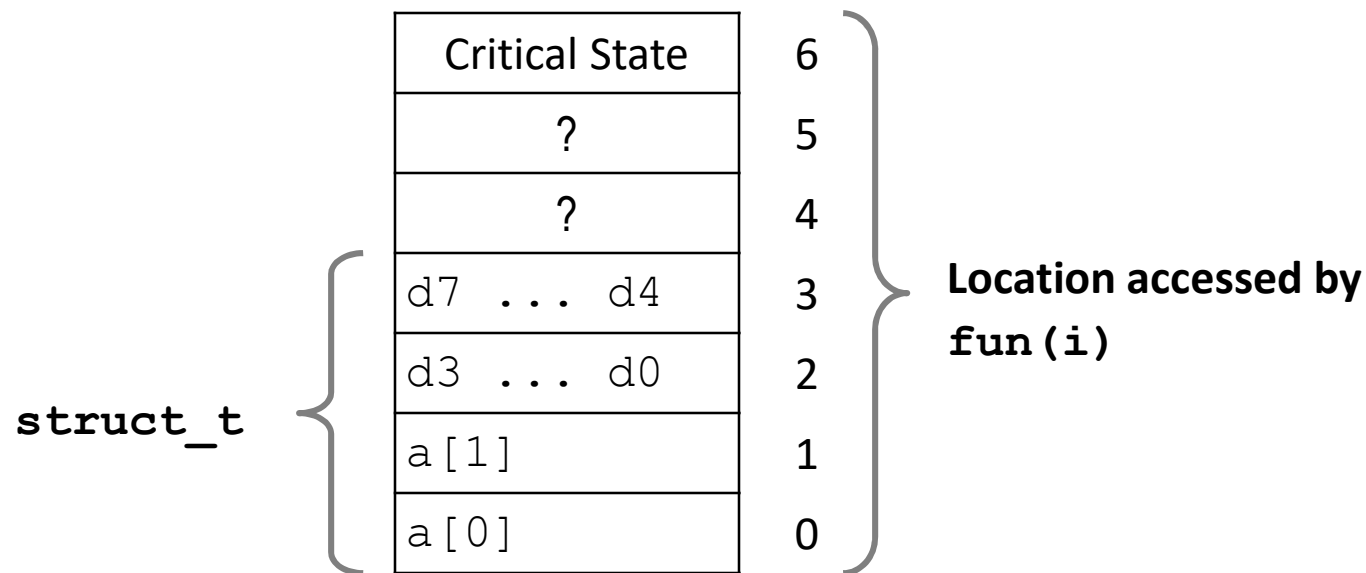


Memory Referencing Bug Example

```
typedef struct {
    int a[2];
    double d;
} struct_t;
```

fun(0)	↯	3.14
fun(1)	↯	3.14
fun(2)	↯	3.1399998664856
fun(3)	↯	2.00000061035156
fun(4)	↯	3.14
fun(6)	↯	Segmentation fault

Explanation:



Such problems are a BIG deal

- **Generally called a “buffer overflow”**

- when exceeding the memory size allocated for an array

- **Why a big deal?**

- It's the #1 technical cause of security vulnerabilities
 - #1 overall cause is social engineering / user ignorance

- **Most common form**

- Unchecked lengths on string inputs
 - Particularly for bounded character arrays on the stack
 - sometimes referred to as stack smashing



String Library Code

Implementation of Unix function `gets()`

```
/* Get string from stdin */
char *gets(char *dest)
{
    int c = getchar();
    char *p = dest;
    while (c != EOF && c != '\n') {
        *p++ = c;
        c = getchar();
    }
    *p = '\0';
    return dest;
}
```

No way to specify limit on number of characters to read

Similar problems with other library functions

`strcpy`, `strcat`: Copy strings of arbitrary length

`scanf`, `fscanf`, `sscanf`, when given `%s` conversion specification



Vulnerable Buffer Code

```
/* Echo Line */  
void echo()  
{  
    char buf[4]; /* Way too small! */  
    gets(buf);  
    puts(buf);  
}
```

```
void call_echo() {  
    echo();  
}
```

← btw, how big
is big enough?

```
unix>./bufdemo-nsp  
Type a string:012345678901234567890123  
012345678901234567890123
```

```
unix>./bufdemo-nsp  
Type a string:0123456789012345678901234  
Segmentation Fault
```



Buffer Overflow Disassembly

echo:

00000000004006cf <echo>:

4006cf:	48 83 ec 18	sub	\$0x18 , %rsp
4006d3:	48 89 e7	mov	%rsp , %rdi
4006d6:	e8 a5 ff ff ff	callq	400680 <gets>
4006db:	48 89 e7	mov	%rsp, %rdi
4006de:	e8 3d fe ff ff	callq	400520 <puts@plt>
4006e3:	48 83 c4 18	add	\$0x18, %rsp
4006e7:	c3	retq	

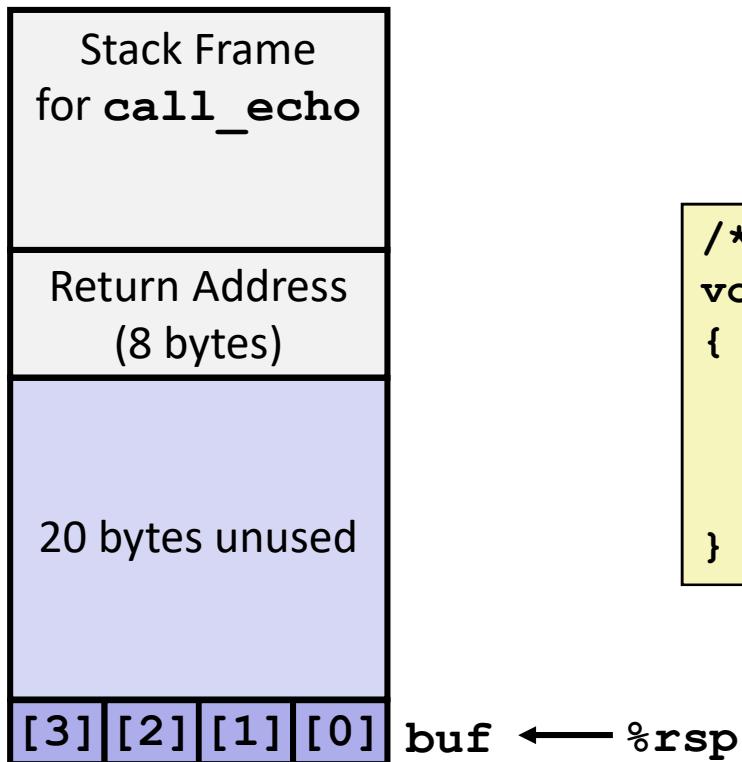
call_echo:

4006e8:	48 83 ec 08	sub	\$0x8, %rsp
4006ec:	b8 00 00 00 00	mov	\$0x0, %eax
4006f1:	e8 d9 ff ff ff	callq	4006cf <echo>
4006f6:	48 83 c4 08	add	\$0x8 , %rsp
4006fa:	c3	retq	



Buffer Overflow Stack

Before call to gets




```
/* Echo Line */  
void echo()  
{  
    char buf[4]; /* Way too small! */  
    gets(buf);  
    puts(buf);  
}
```

```
echo:  
    subq    $24, %rsp  
    movq    %rsp, %rdi  
    call    gets  
    . . .
```

Buffer Overflow Stack Example

Before call to gets



```
void echo()  
{  
    char buf[4];  
    gets(buf);  
    . . .  
}
```

```
echo:  
    subq    $24, %rsp  
    movq    %rsp, %rdi  
    call    gets  
    . . .
```

call_echo:

```
. . .  
4006f1:    callq   4006cf <echo>  
4006f6:    add     $0x8,%rsp  
. . .
```



Buffer Overflow Stack Example #1

After call to gets

Stack Frame for <code>call_echo</code>			
00	00	00	00
00	40	06	f6
00	32	31	30
39	38	37	36
35	34	33	32
31	30	39	38
37	36	35	34
33	32	31	30

```
void echo()  
{  
    char buf[4];  
    gets(buf);  
    . . .  
}
```

```
echo:  
    subq    $24, %rsp  
    movq    %rsp, %rdi  
    call    gets  
    . . .
```

`call_echo:`

```
. . .  
4006f1:  callq    4006cf <echo>  
4006f6:  add      $0x8, %rsp  
. . .
```

`buf` ← `%rsp`

```
unix> ./bufdemo-nsp  
Type a string: 01234567890123456789012  
01234567890123456789012
```

Overflowed buffer, but did not corrupt state



Buffer Overflow Stack Example #2

After call to gets

Stack Frame for <code>call_echo</code>			
00	00	00	00
00	40	00	34
33	32	31	30
39	38	37	36
35	34	33	32
31	30	29	28
27	26	25	24
23	22	21	20

```
void echo()  
{  
    char buf[4];  
    gets(buf);  
    . . .  
}
```

```
echo:  
    subq    $24, %rsp  
    movq    %rsp, %rdi  
    call    gets  
    . . .
```

`call_echo:`

```
. . .  
4006f1:    callq   4006cf <echo>  
4006f6:    add     $0x8,%rsp  
. . .
```

`buf ← %rsp`

```
unix> ./bufdemo-nsp  
Type a string: 0123456789012345678901234  
Segmentation Fault
```

Overflowed buffer and corrupted return pointer



Buffer Overflow Stack Example #3

After call to gets

Stack Frame for <code>call_echo</code>			
00	00	00	00
00	40	06	00
33	32	31	30
39	38	37	36
35	34	33	32
31	30	39	38
37	36	35	34
33	32	31	30

```
void echo()  
{  
    char buf[4];  
    gets(buf);  
    . . .  
}
```

```
echo:  
    subq    $24, %rsp  
    movq    %rsp, %rdi  
    call    gets  
    . . .
```

`call_echo:`

```
. . .  
4006f1:    callq   4006cf <echo>  
4006f6:    add     $0x8,%rsp  
. . .
```

`buf ← %rsp`

```
unix> ./bufdemo-nsp  
Type a string: 012345678901234567890123  
012345678901234567890123
```

Overflowed buffer, corrupted return pointer, but program seems to work!



Buffer Overflow Stack Example #3 Explained

After call to gets

Stack Frame for <code>call_echo</code>			
00	00	00	00
00	40	06	00
33	32	31	30
39	38	37	36
35	34	33	32
31	30	39	38
37	36	35	34
33	32	31	30

`buf` ← `%rsp`

`register_tm_clones:`

```
. . .
400600:  mov    %rsp,%rbp
400603:  mov    %rax,%rdx
400606:  shr    $0x3f,%rdx
40060a:  add    %rdx,%rax
40060d:  sar    %rax
400610:  jne    400614
400612:  pop    %rbp
400613:  retq
```

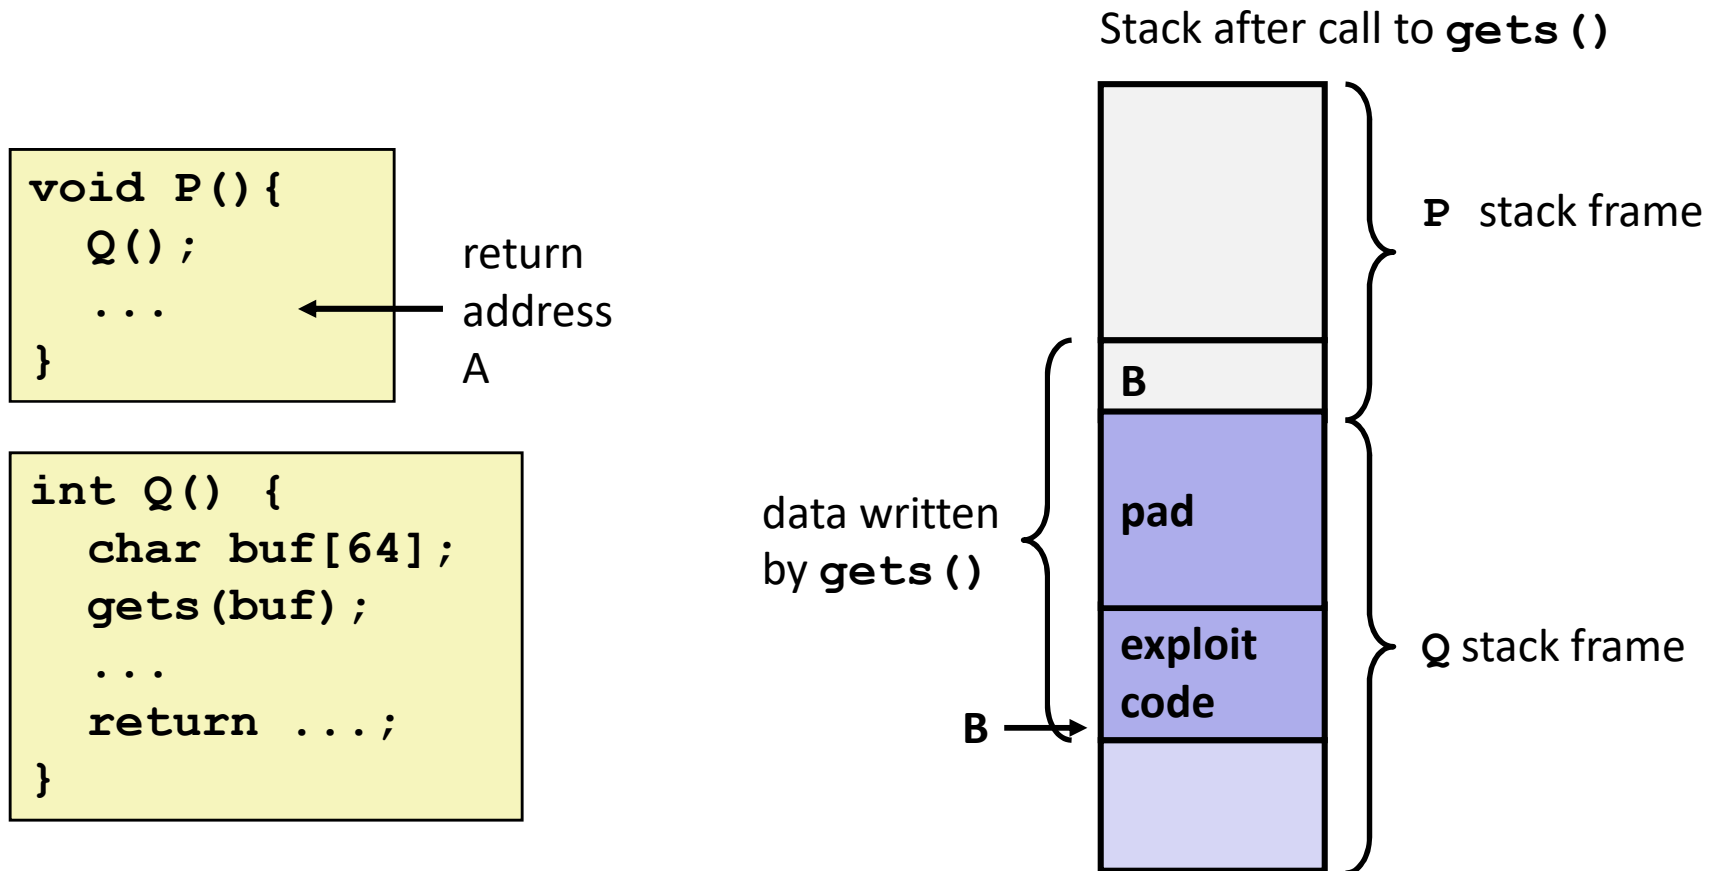
“Returns” to unrelated code

Lots of things happen, without modifying critical state

Eventually executes `retq` back to main



Code Injection Attacks



- Input string contains byte representation of executable code
- Overwrite return address A with address of buffer B
- When Q executes `ret`, will jump to exploit code



Exploits Based on Buffer Overflows

- 🌀 *Buffer overflow bugs can allow remote machines to execute arbitrary code on victim machines*
- 🌀 **Distressingly common in real programs**
 - 🌀 Programmers keep making the same mistakes ☹️
 - 🌀 Recent measures make these attacks much more difficult
- 🌀 **Examples across the decades**
 - 🌀 Original “Internet worm” (1988)
 - 🌀 “IM wars” (1999)
 - 🌀 Twilight hack on Wii (2000s)
 - 🌀 ... and many, many more
- 🌀 **You will learn some of the tricks in attacklab**
 - 🌀 Hopefully to convince you to never leave such holes in your programs!!



Example: the original Internet worm (1988)

🌀 Exploited a few vulnerabilities to spread

- 🌀 Early versions of the finger server (fingerd) used `gets ()` to read the argument sent by the client:

- 🌀 `finger droh@cs.cmu.edu`

- 🌀 Worm attacked fingerd server by sending phony argument:

- 🌀 `finger "exploit-code padding new-return-address"`

- 🌀 exploit code: executed a root shell on the victim machine with a direct TCP connection to the attacker.

🌀 Once on a machine, scanned for other machines to attack

- 🌀 invaded ~6000 computers in hours (10% of the Internet 😊)

- 🌀 see June 1989 article in *Comm. of the ACM*

- 🌀 the young author of the worm was prosecuted...


- 🌀 and CERT was formed... still homed at CMU



Example 2: IM War

🕒 July, 1999

- 🕒 Microsoft launches MSN Messenger (instant messaging system).
- 🕒 Messenger clients can access popular AOL Instant Messaging Service (AIM) servers



IM War (cont.)

🌀 August 1999

- 🌀 Mysteriously, Messenger clients can no longer access AIM servers
- 🌀 Microsoft and AOL begin the IM war:
 - 🌀 AOL changes server to disallow Messenger clients
 - 🌀 Microsoft makes changes to clients to defeat AOL changes
 - 🌀 At least 13 such skirmishes
- 🌀 What was really happening?
 - 🌀 AOL had discovered a buffer overflow bug in their own AIM clients
 - 🌀 They exploited it to detect and block Microsoft: the exploit code returned a 4-byte signature (the bytes at some location in the AIM client) to server
 - 🌀 When Microsoft changed code to match signature, AOL changed signature location



OK, what to do about buffer overflow attacks

- ➊ Avoid overflow vulnerabilities
- ➋ Employ system-level protections
- ➌ Have compiler use “stack canaries”
- ➍ Lets talk about each...



1. Avoid Overflow Vulnerabilities in Code (!)

```
/* Echo Line */  
void echo()  
{  
    char buf[4]; /* Way too small! */  
    fgets(buf, 4, stdin);  
    puts(buf);  
}
```

🌀 For example, use library routines that limit string lengths

- 🌀 **fgets** instead of **gets**
- 🌀 **strncpy** instead of **strcpy**
- 🌀 Don't use **scanf** with **%s** conversion specification
 - 🌀 Use **fgets** to read the string
 - 🌀 Or use **%ns** where **n** is a suitable integer




2. System-Level Protections can help

🌀 Randomized stack offsets

- 🌀 At start of program, allocate random amount of space on stack
- 🌀 Shifts stack addresses for entire program
- 🌀 Makes it difficult for hacker to predict beginning of inserted code
- 🌀 E.g.: 5 executions of memory allocation code

local 0x7ffe4d3be87c 0x7fff75a4f9fc 0x7ffeadb7c80c 0x7ffeaea2fdac 0x7ffcd452017c

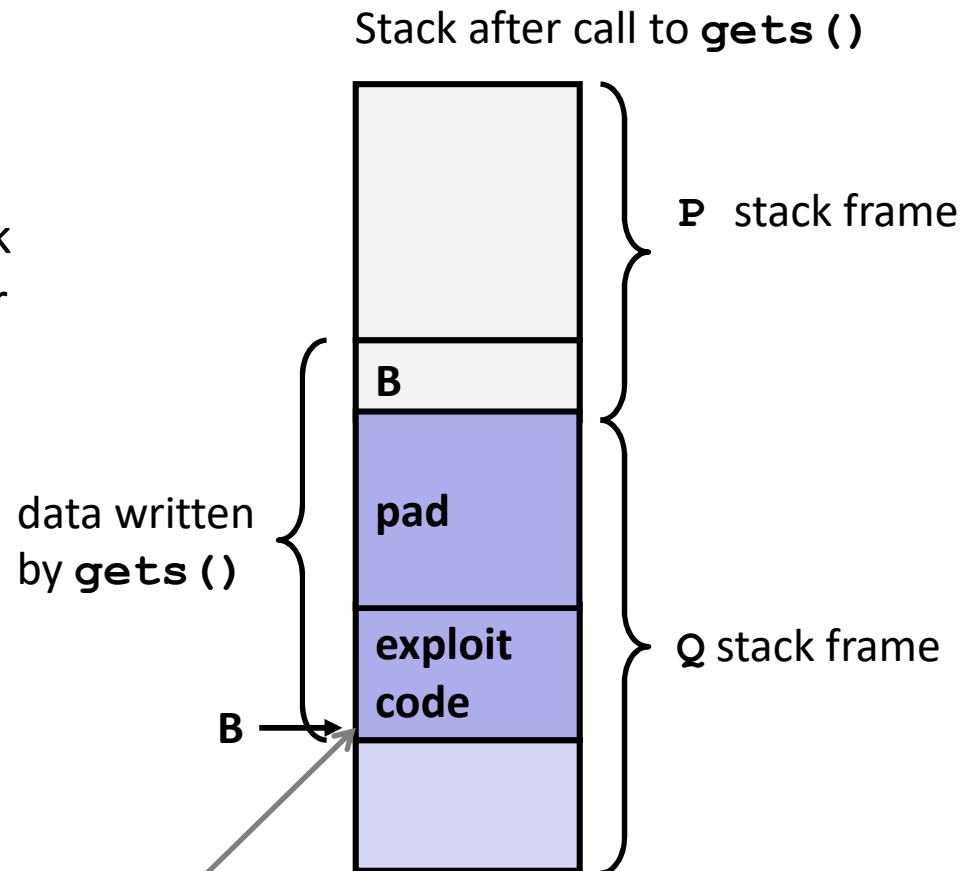
- 🌀 Stack repositioned each time program executes



2. System-Level Protections can help

Nonexecutable code segments

- In traditional x86, can mark region of memory as either “read-only” or “writeable”
 - Can execute anything readable
- X86-64 added explicit “execute” permission
- Stack marked as non-executable



Any attempt to execute this code will fail



3. Stack Canaries can help

🌀 Idea

- 🌀 Place special value (“canary”) on stack just beyond buffer
- 🌀 Check for corruption before exiting function

🌀 GCC Implementation

- 🌀 `-fstack-protector`
- 🌀 Now the default (disabled earlier)

```
unix> ./bufdemo-sp  
Type a string: 0123456  
0123456
```

```
unix> ./bufdemo-sp  
Type a string: 01234567  
*** stack smashing detected ***
```



Protected Buffer Disassembly

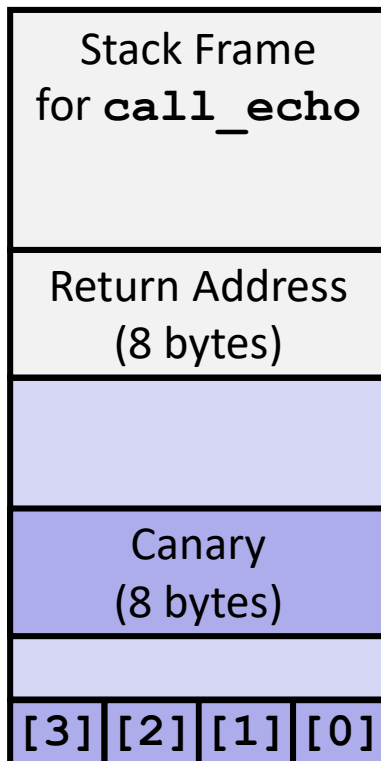
echo:

```
40072f:  sub    $0x18,%rsp
400733:  mov    %fs:0x28,%rax
40073c:  mov    %rax,0x8(%rsp)
400741:  xor    %eax,%eax
400743:  mov    %rsp,%rdi
400746:  callq  4006e0 <gets>
40074b:  mov    %rsp,%rdi
40074e:  callq  400570 <puts@plt>
400753:  mov    0x8(%rsp),%rax
400758:  xor    %fs:0x28,%rax
400761:  je     400768 <echo+0x39>
400763:  callq  400580 <__stack_chk_fail@plt>
400768:  add    $0x18,%rsp
40076c:  retq
```



Setting Up Canary

Before call to gets



```
/* Echo Line */  
void echo()  
{  
    char buf[4]; /* Way too small! */  
    gets(buf);  
    puts(buf);  
}
```

```
echo:  
    . . .  
    movq    %fs:40, %rax    # Get canary  
    movq    %rax, 8(%rsp)  # Place on stack  
    xorl    %eax, %eax      # Erase canary  
    . . .
```



Checking Canary

After call to gets

Stack Frame for <code>call_echo</code>			
Return Address (8 bytes)			
Canary (8 bytes)			
00	36	35	34
33	32	31	30

```
/* Echo Line */
void echo()
{
    char buf[4]; /* Way too small! */
    gets(buf);
    puts(buf);
}
```

Input: 0123456

buf ← %rsp

```
echo:
    . . .
    movq    8(%rsp), %rax    # Retrieve from
stack
    xorq    %fs:40, %rax    # Compare to canary
    je      .L6             # If same, OK
    call    __stack_chk_fail # FAIL
.L6:
```



Return-Oriented Programming Attacks

⌚ Challenge (for hackers)

- ⌚ Stack randomization makes it hard to predict buffer location
- ⌚ Marking stack nonexecutable makes it hard to insert binary code

⌚ Alternative Strategy

- ⌚ Use existing code
 - ⌚ E.g., library code from `stdlib`
- ⌚ String together fragments to achieve overall desired outcome
- ⌚ *Does not overcome stack canaries*

⌚ Construct program from *gadgets*

- ⌚ Sequence of instructions ending in `ret`
 - ⌚ Encoded by single byte `0xc3`
- ⌚ Code positions fixed from run to run
- ⌚ Code is executable



Gadget Example #1

```
long ab_plus_c
(long a, long b, long c)
{
    return a*b + c;
}
```

```
00000000004004d0 <ab_plus_c>:
4004d0: 48 0f af fe  imul %rsi,%rdi
4004d4: 48 8d 04 17  lea (%rdi,%rdx,1),%rax
4004d8: c3           retq
```

$\text{rax} \leftarrow \text{rdi} + \text{rdx}$

Gadget address = 0x4004d4

 Use tail end of existing functions



Gadget Example #2

```
void setval(unsigned *p) {  
    *p = 3347663060u;  
}
```

Encodes `movq %rax, %rdi`

<setval>:							
4004d9:	c7	07	d4	48	89	c7	<code>movl \$0xc78948d4, (%rdi)</code>
4004df:	c3						<code>retq</code>


$\text{rdi} \leftarrow \text{rax}$

Gadget address = 0x4004dc

 Repurpose byte codes



ROP Execution



- Trigger with `ret` instruction
 - Will start executing Gadget 1
- Final `ret` in each gadget will start next one



Today

- Memory Layout

- Buffer Overflow

 - Vulnerability

 - Protection

- Unions



Union Allocation

- Allocate according to largest element
- Can only use one field at a time

```
union U1 {  
    char c;  
    int i[2];  
    double v;  
} *up;
```

```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *sp;
```

